

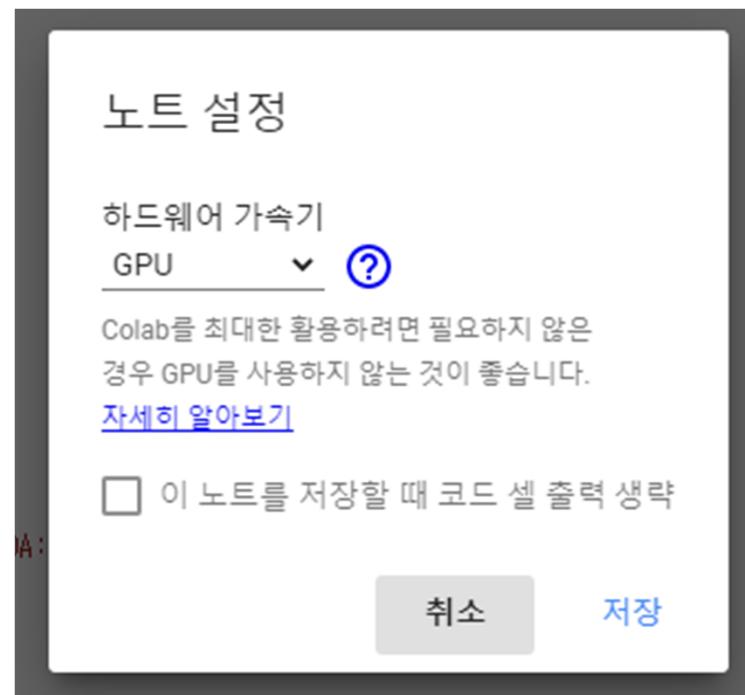
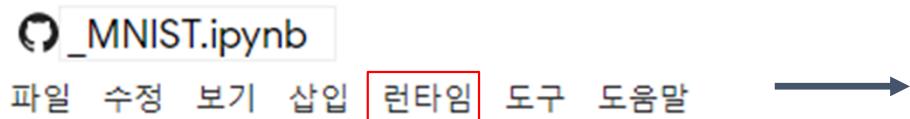
# Week 3: MNIST using MLP

KAIST

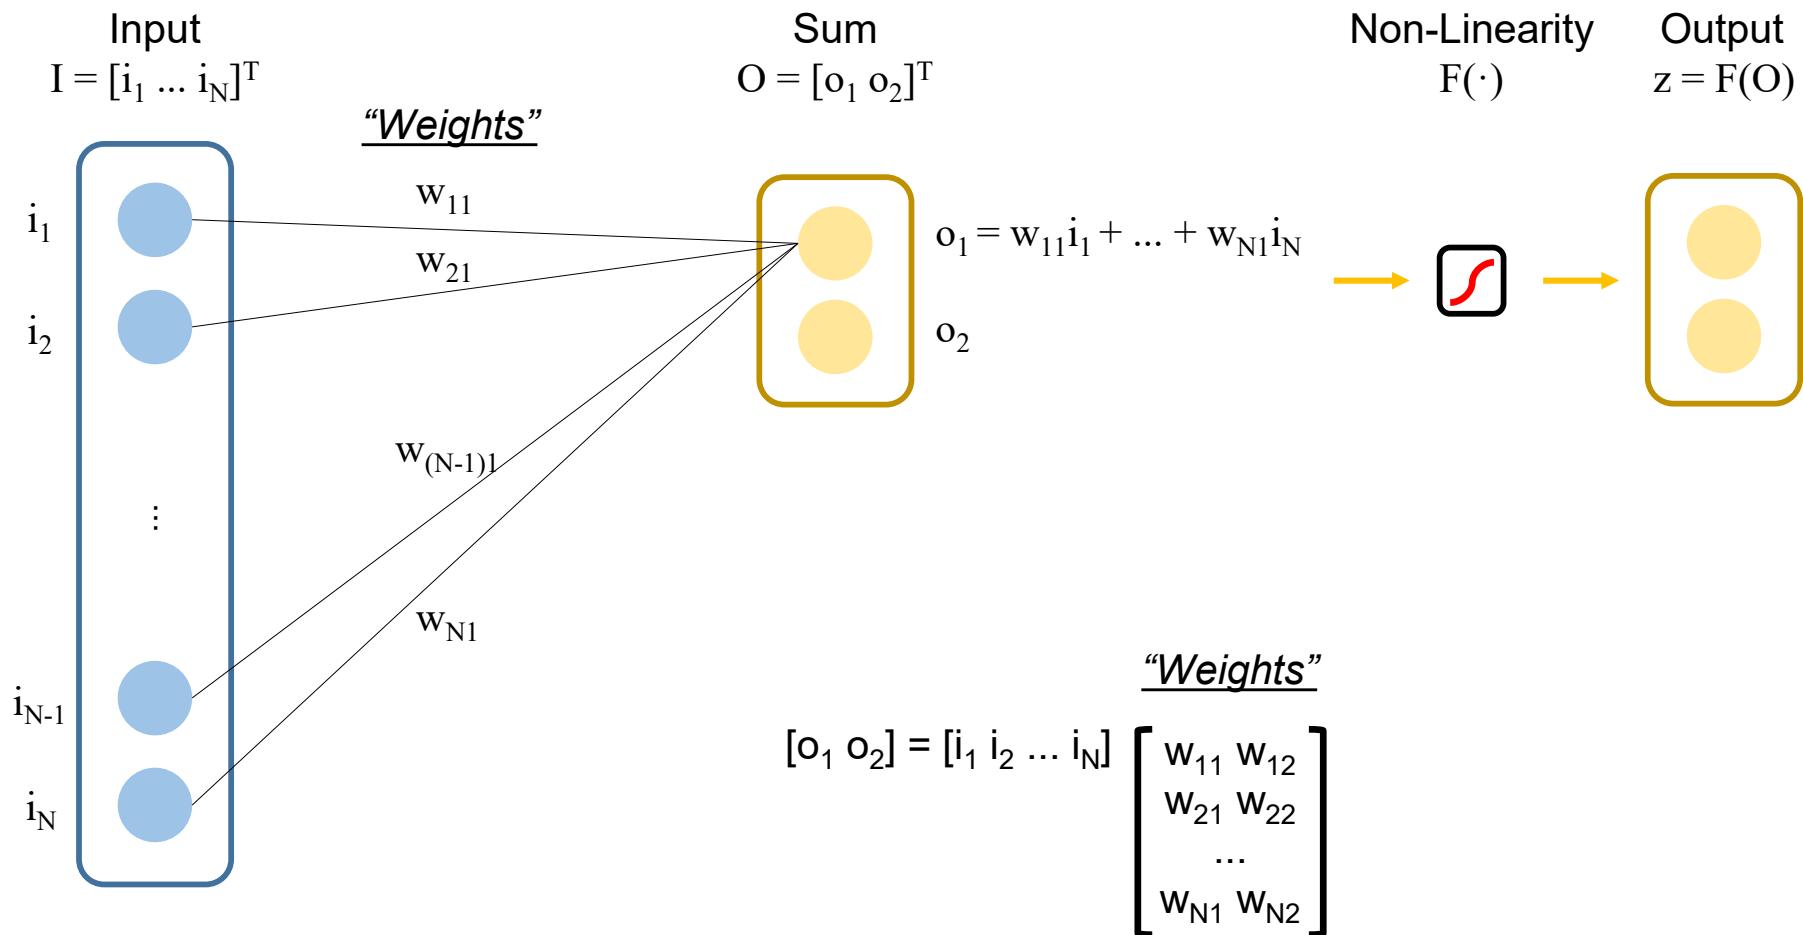
Mobile Robotics & Intelligence Lab

박사과정 김진식

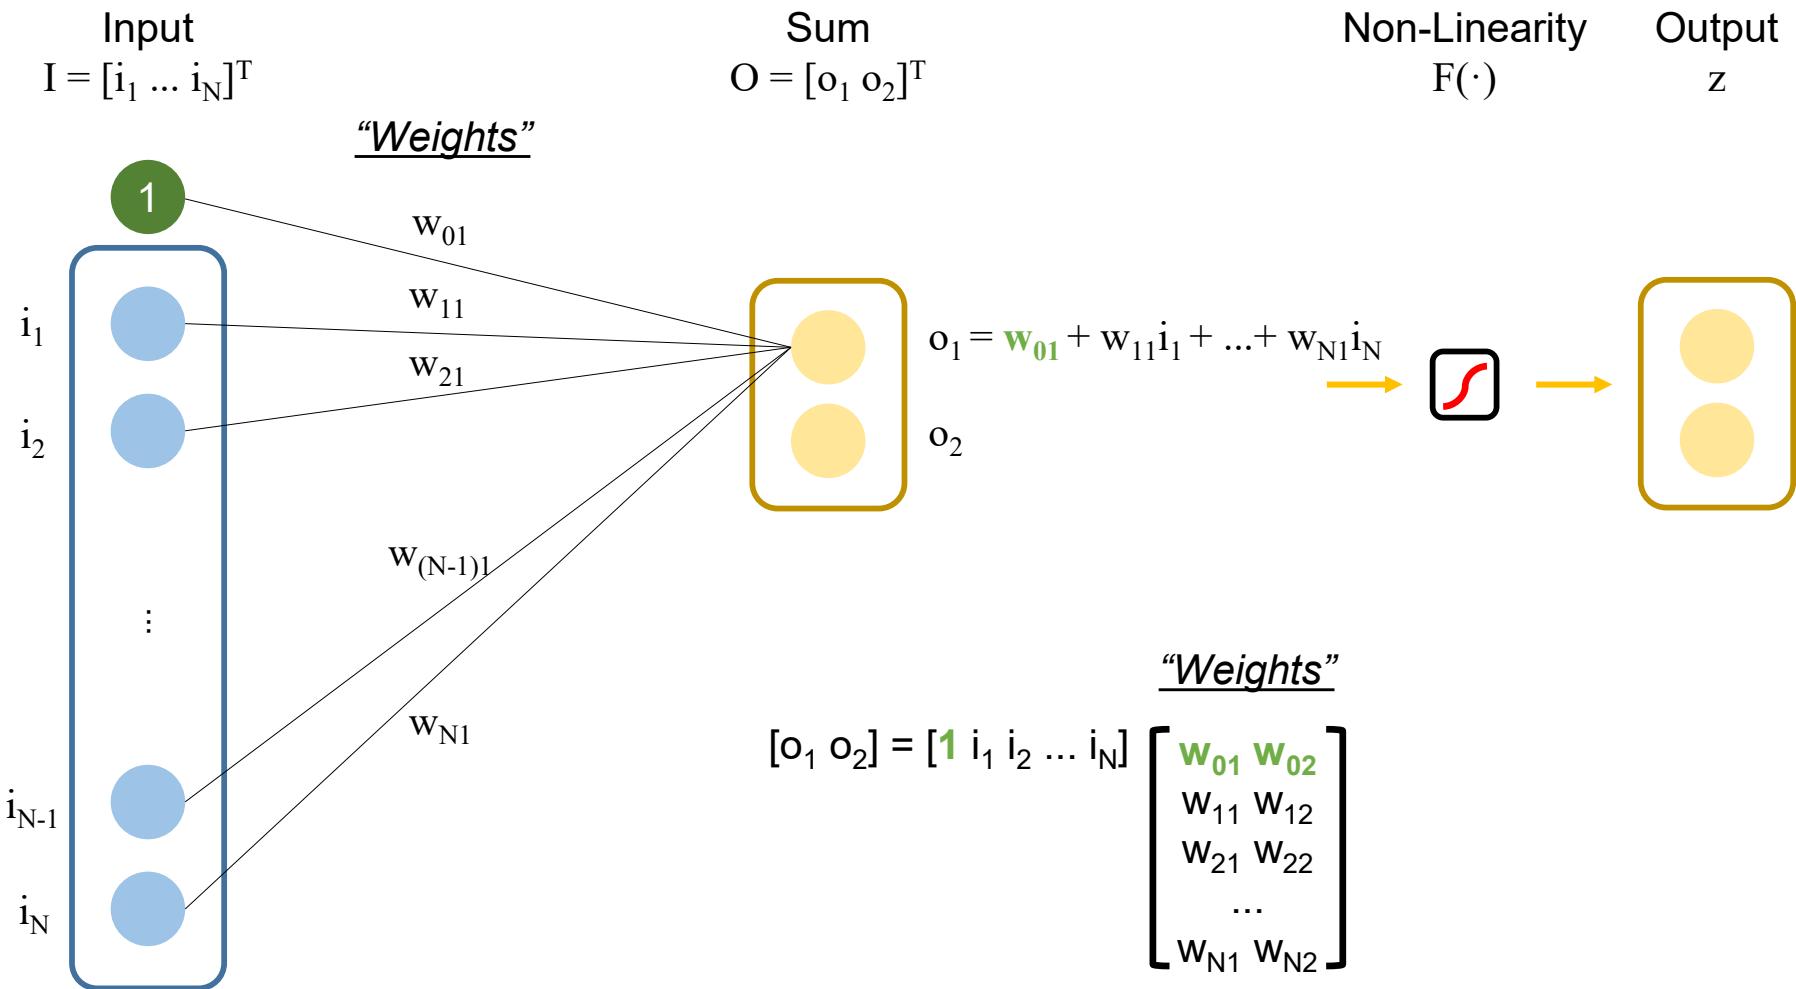
런타임 -> 런타임 유형 변경 -> 하드웨어 가속  
기(GPU)



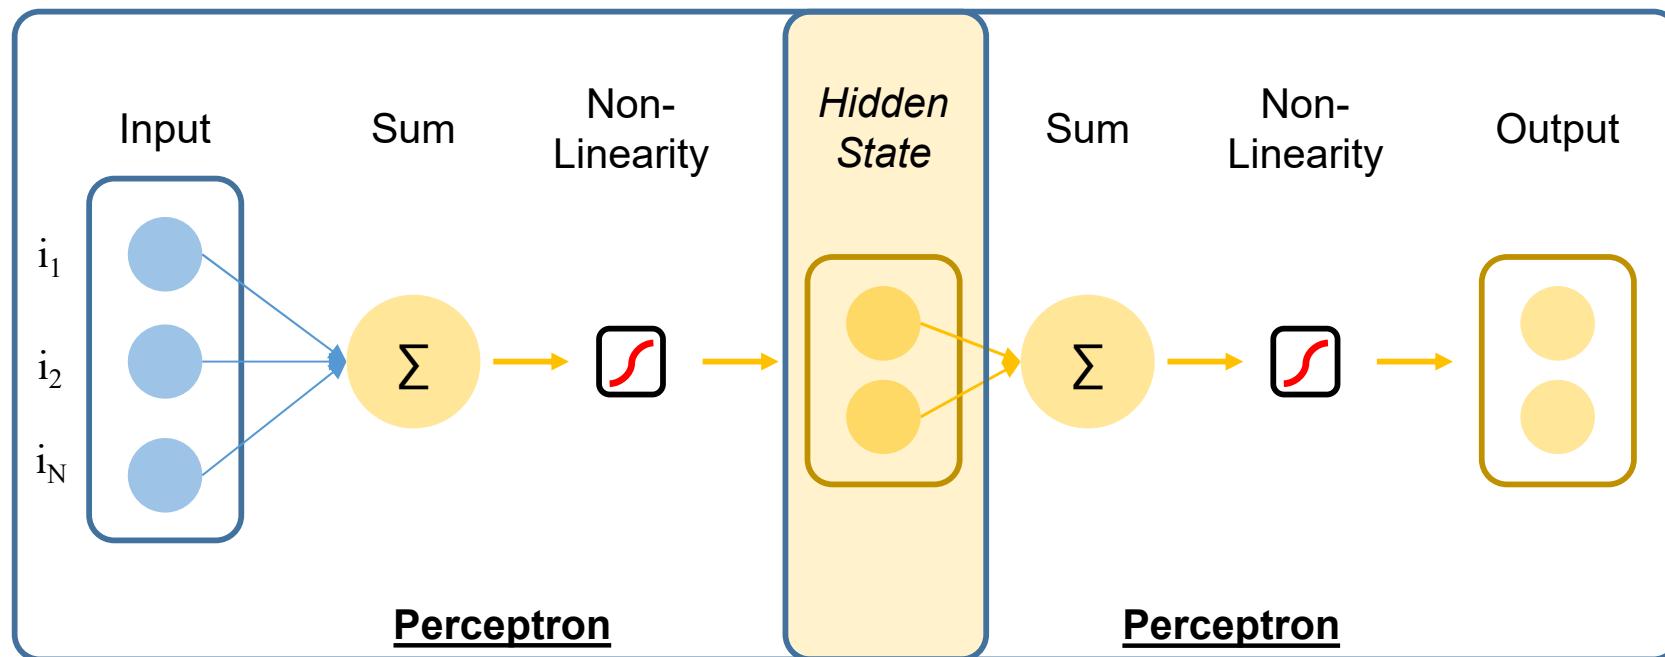
## Perceptrons (No bias)



## Perceptrons (with bias)



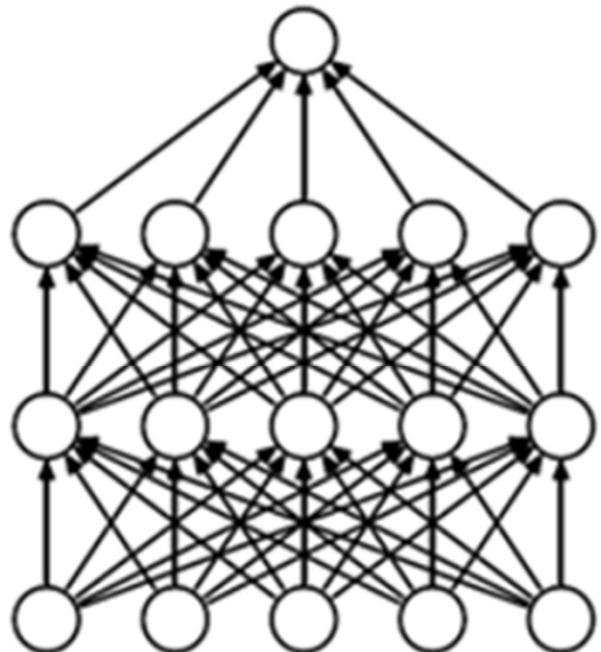
## Multi-layer perceptron



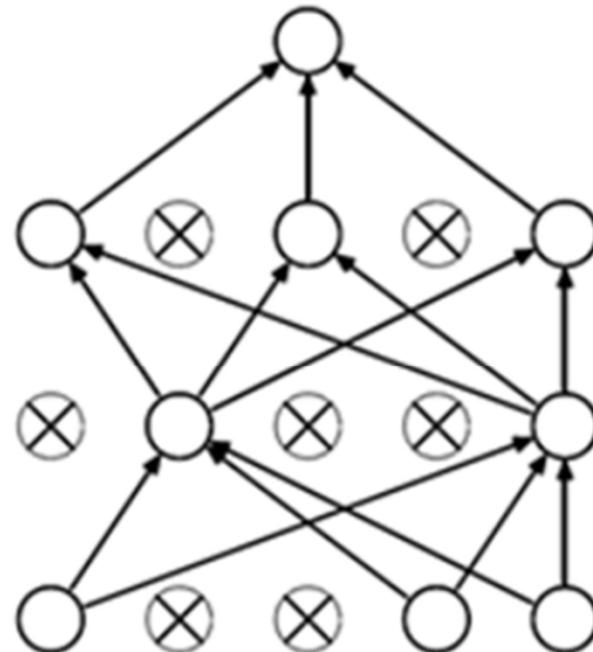
+ Dropout

## Multi-layer perceptron – dropout

- One of the biggest problems of deep learning is overfitting to training data.
- *Dropout rate* defines probability of each neuron to be *zeroed out*.
- We can partially solve the overfitting problem through dropout method.



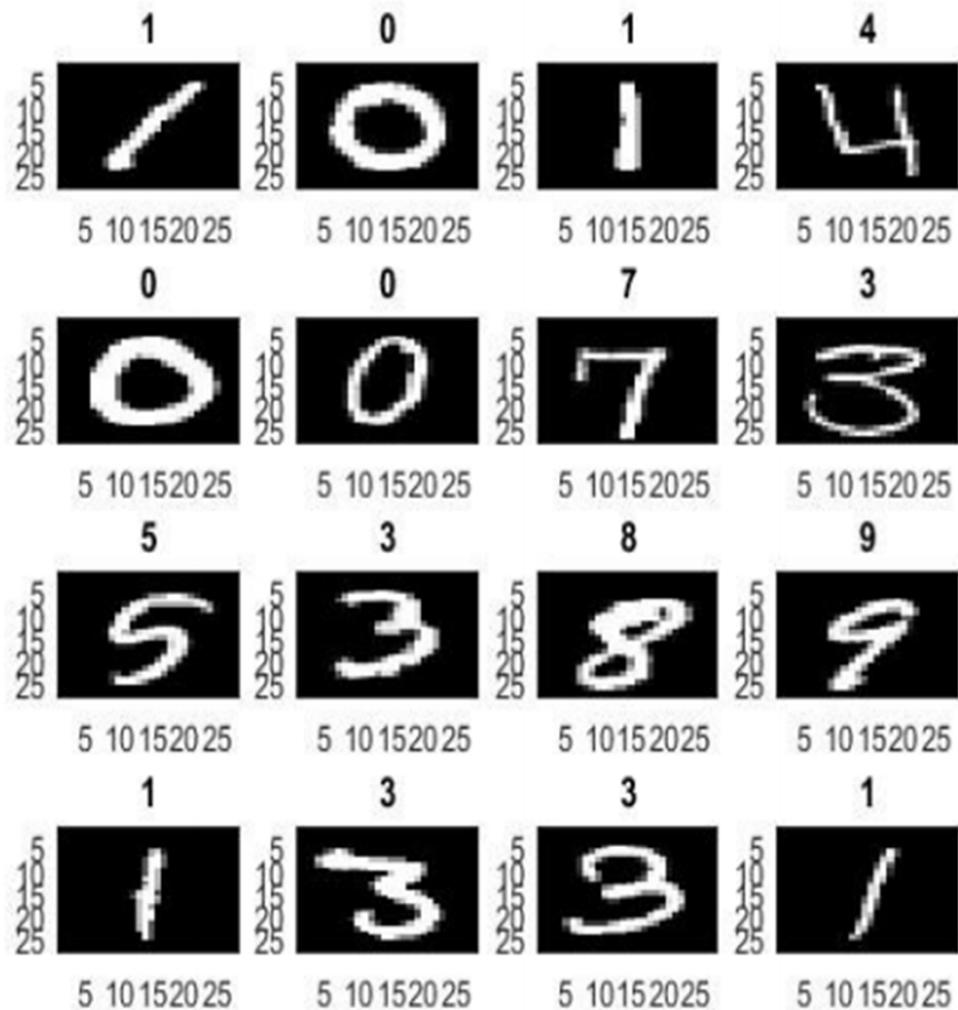
(a) Standard Neural Net



(b) After applying dropout.

## MNIST Dataset

- MNIST: Large hand written digit classification database
- Format
  - Input: 28 x 28 gray scale image
  - Output: 10 labels(0-9)
  - Centered on center of the mass



**Today :** Image( $28 \times 28$ ) classification using MLP

## MNIST tutorial

### Dataset preparation

```
batch_size = 32

kwargs = {'num_workers': 1, 'pin_memory': True} if cuda else {}

# This will automatically download MNIST data into ../data directory
train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=True, download=True,
                  transform=transforms.Compose([
                      transforms.ToTensor(),
                      transforms.Normalize((0.1307,), (0.3081,))
                  ])),
    batch_size=batch_size, shuffle=True, **kwargs)

validation_loader = torch.utils.data.DataLoader(
    datasets.MNIST('../data', train=False, transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))
    ])),
    batch_size=batch_size, shuffle=False, **kwargs)
```

Training

Validation

- Load MNIST dataset using data loader
- Input size: 28 x 28 gray scale image
- Output: classes of ("0", .. "9") for each training digit

## MNIST tutorial

### Dataloader usage

```
for (X_train, y_train) in train_loader:  
    print('X_train:', X_train.size(), 'type:', X_train.type())  
    print('y_train:', y_train.size(), 'type:', y_train.type())  
    break
```

X\_train: torch.Size([32, 1, 28, 28]) type: torch.FloatTensor  
y\_train: torch.Size([32]) type: torch.LongTensor

X\_train (torch.FloatTensor) : Image



Batch Size (=32)

Y\_train (torch.LongTensor) : Label

6, 6, 5, 7, 4, ...

Batch Size

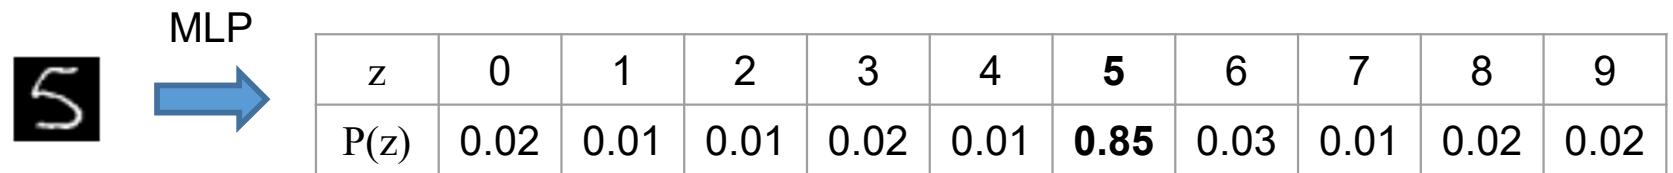
- “Batch” : Divide dataset into smaller sets

X\_train.shape : (Batch Size) × (Channels #) × (Width) × (Height)      #  $32 \times 1 \times 28 \times 28$   
Y\_train.shape : (Batch Size)    # 32

## MNIST classification with MLP



### Example



The example shows a handwritten digit '5' being processed by an MLP. The output is a probability distribution over the digits 0 through 9, represented in a table:

z	0	1	2	3	4	5	6	7	8	9
P(z)	0.02	0.01	0.01	0.02	0.01	<b>0.85</b>	0.03	0.01	0.02	0.02

- Softmax

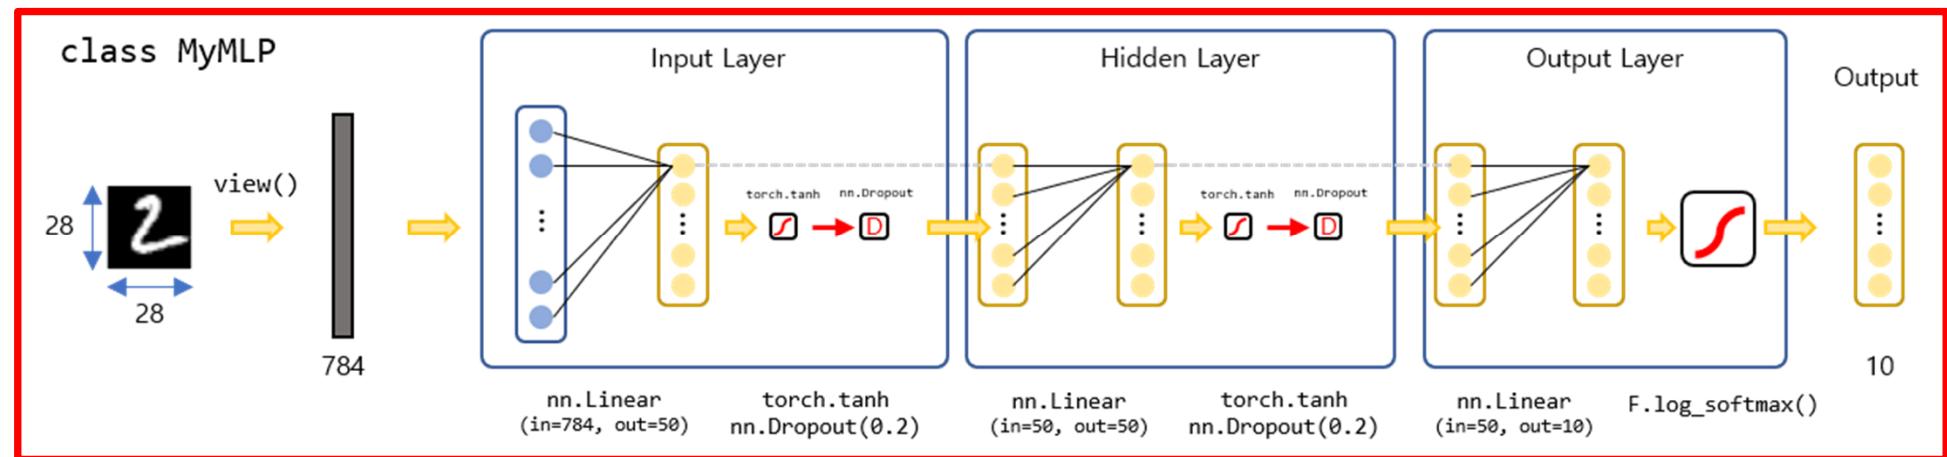
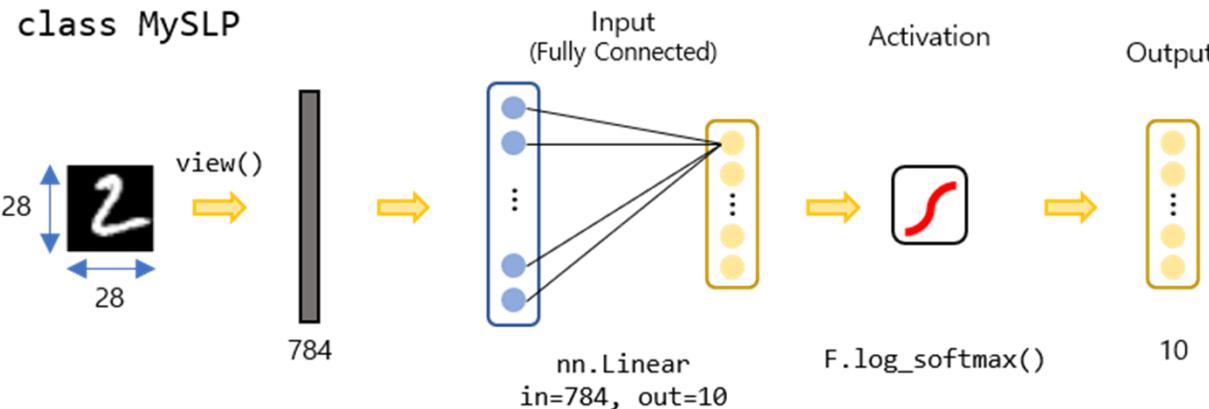
$$\mathbf{z} = [z_1, \dots, z_n]$$
$$\text{softmax}(\mathbf{z}) = [e^{z_1}/\sum(e^{z_k}), e^{z_2}/\sum(e^{z_k}), \dots, e^{z_n}/\sum(e^{z_k})]$$

**Today, you will be asked to**

- Task 1. Implement custom MLP class
- Task 2. Complete training pipeline  
+ Visualization

**Check out colab notebook for details**

## Implement MyMLP class



## Pytorch MLP Basic – nn.Linear

CLASS `torch.nn.Linear(in_features: int, out_features: int, bias: bool = True)`

[SOURCE]

Applies a linear transformation to the incoming data:  $y = xA^T + b$

### Parameters

- **in\_features** – size of each input sample
- **out\_features** – size of each output sample
- **bias** – If set to `False`, the layer will not learn an additive bias. Default: `True`

### Shape:

- Input:  $(N, *, H_{in})$  where  $*$  means any number of additional dimensions and  $H_{in} = \text{in\_features}$
- Output:  $(N, *, H_{out})$  where all but the last dimension are the same shape as the input and  $H_{out} = \text{out\_features}$ .

### Variables

- **~Linear.weight** – the learnable weights of the module of shape  $(\text{out\_features}, \text{in\_features})$ . The values are initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$ , where  $k = \frac{1}{\text{in\_features}}$
- **~Linear.bias** – the learnable bias of the module of shape  $(\text{out\_features})$ . If `bias` is `True`, the values are initialized from  $\mathcal{U}(-\sqrt{k}, \sqrt{k})$  where  $k = \frac{1}{\text{in\_features}}$

<https://pytorch.org/docs/master/generated/torch.nn.Linear.html>

## Pytorch MLP Basic – nn.Dropout

CLASS `torch.nn.Dropout(p: float = 0.5, inplace: bool = False)`

[SOURCE]

During training, randomly zeroes some of the elements of the input tensor with probability `p` using samples from a Bernoulli distribution. Each channel will be zeroed out independently on every forward call.

This has proven to be an effective technique for regularization and preventing the co-adaptation of neurons as described in the paper [Improving neural networks by preventing co-adaptation of feature detectors](#).

Furthermore, the outputs are scaled by a factor of  $\frac{1}{1-p}$  during training. This means that during evaluation the module simply computes an identity function.

### Parameters

- `p` – probability of an element to be zeroed. Default: 0.5
- `inplace` – If set to `True`, will do this operation in-place. Default: `False`

### Shape:

- Input: `(*)`. Input can be of any shape
- Output: `(*)`. Output is of the same shape as input

<https://pytorch.org/docs/master/generated/torch.nn.Dropout.html>

## Complete training pipeline (2-1)

### Training

```
def train(model, optimizer, epoch, log_interval=100):
    model.train()    # Set model to training mode

    for batch_idx, (data, target) in enumerate(train_loader):
        if cuda:
            data, target = data.cuda(), target.cuda()
        data, target = Variable(data), Variable(target)

        # 2-1. Clear remaining gradients by setting to zero
        # 2-2. Call forward() of model
        # 2-3. Compute loss (use F.nll_loss)
        # 2-4. Compute gradients
        # 2-5. Backpropagation by updating weights
        pass

## CLI Visualization (Optional) ##
    if ((batch_idx % log_interval) == 0) or (batch_idx == len(train_loader)-1):
```

## Complete training pipeline (2-1)

### Training

```
def train(epoch, log_interval=100):
    model.train() → Set to train mode
    for batch_idx, (data, target) in enumerate(train_loader):
        if cuda:
            data, target = data.cuda(), target.cuda()
        data, target = Variable(data), Variable(target)
        optimizer.zero_grad()
        output = model(data) → Get output
        loss = F.nll_loss(output, target) → Calculate loss
        loss.backward() → Backpropagation using optimizer
        optimizer.step()
        if batch_idx % log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.data[0])) → Print training log
```

## Validation (Complete)

```
def validate(loss_vector, accuracy_vector):
    model.eval() → Set to validation mode
    val_loss, correct = 0, 0
    for data, target in validation_loader:
        if cuda:
            data, target = data.cuda(), target.cuda()
        data, target = Variable(data, volatile=True), Variable(target)
        output = model(data)
        val_loss += F.nll_loss(output, target).data[0]
        pred = output.data.max(1)[1] # get the index of the max log-probability
        correct += pred.eq(target.data).cpu().sum()

    val_loss /= len(validation_loader) } → Calculate validation loss
    loss_vector.append(val_loss)

    accuracy = 100. * correct / len(validation_loader.dataset) } → Calculate accuracy of validation set
    accuracy_vector.append(accuracy)

    print('Validation set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)'.format(
        val_loss, correct, len(validation_loader.dataset), accuracy))
```

↓

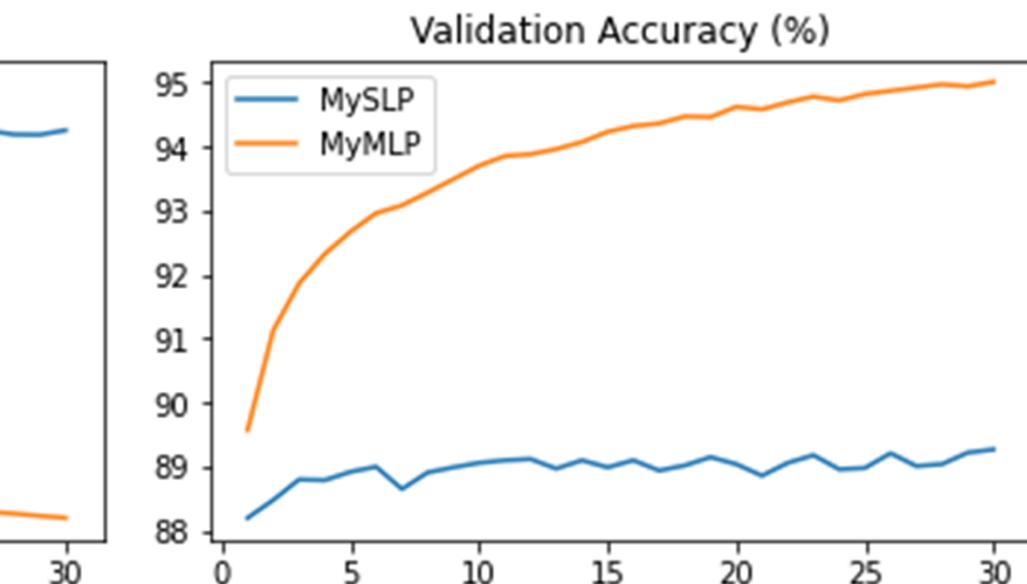
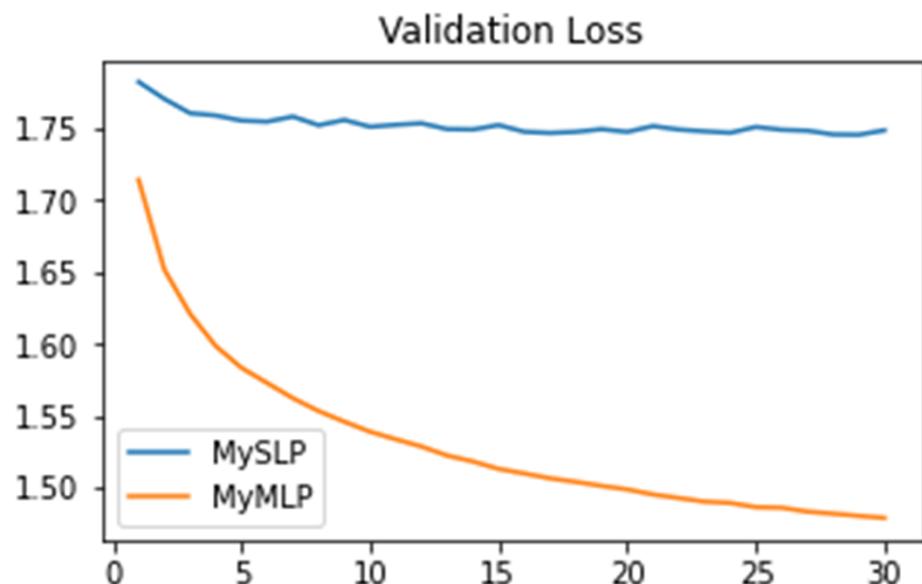
Print log of accuracy of validation set

## Visualization (2-3)

```
plt.figure(figsize=(5,3))
plt.plot(np.arange(1,epochs+1), lossv)
plt.title('validation loss')

plt.figure(figsize=(5,3))
plt.plot(np.arange(1,epochs+1), accv)
plt.title('validation accuracy');
```

→ **Plot the loss and accuracy graph  
We can visualize the performance and  
the loss of our network in validation  
sets.**



- Today's quiz #1. Complete ~2.3 and submit your outputs (must run by yourself)

## Practice Session / Q & A

Task 1. Implement custom MLP class (1)

Task 2. Complete training pipeline + Visualization (2-1)

- Quiz will begin at 5:30